

DLILO RAG System Technical Report (Summary)

Architecture: ingestion -> chunking -> embeddings -> FAISS -> retrieval -> generation.

Chunking comparison: sentence chunks improved precision@5 vs fixed (see experiment_log.csv).

Evaluation: precision@5, answer relevance, faithfulness (run eval/run_eval.py).

GPT-2 vs BERT: BERT-style encoders for retrieval, GPT-style decoders for generation.

Limitations: small demo corpus, heuristic evaluation, dependence on external LLMs.

Reflection: expand corpus, add RAGAS, improve guardrails and caching.